



Beacon Night Shift

The on-call agent that fixes the 3 AM page with your voice —
and does not wake you the second time.

WHAT ON-CALL ACTUALLY COSTS

53_{min}

median time to recovery, industry-wide

12 % better in five years — against 3× the monitoring spend

74%

of DevOps engineers report burnout

on-call load is the leading indicator; seniors leave first

\$5.6k

per minute of downtime, Gartner baseline

and on a team of four there is nobody to escalate to

WHO THIS IS FOR

**A backend engineer on a team of one to five,
running production on AWS for users in another
time zone.**

No follow-the-sun rotation.
There is no second shift to
escalate to.

A small, repeating failure
surface: a rule dropped by a
deploy, a service that needs a
redploy.

No budget for an SRE
platform, and nobody to run
one.

LIVE, NOT A MOCK-UP

Speak over the read-back, and the fix is withdrawn.



Beacon is speaking — talk to interrupt
heard: "No, wait, stop."
☐ noisy room
Interrupt


BEACON A security group rule that the golden snapshot says should exist is missing, so the
00:15:24 service can no longer —
● INTERRUPTED

A real run on the deployed account: the read-back cut mid-sentence, **INTERRUPTED**, `propose_fix · dry run PASSED`, and the console note — **"Fix 1 withdrawn — you spoke over the read-back, so nothing was applied."** The approval phrase stops working until it is proposed again.

Nine tools on one socket. Consent is decided in code.

Beacon Night Shift — the voice path

Nine tools on one socket. The model decides whether to call one; code decides whether the call is allowed.



The Voice Agent API is the mouth and the ears; every **tool.call** returns to the browser, which runs it on our Lambda with the transcript the API produced.

WHY THIS NEEDED **THIS** GENERATION OF MODELS

- 01 Consent has to rest on **words an engineer actually said** — half asleep, in two languages. A keypad cannot say **which** fix on **which** resource.
- 02 An interruption has to be an **event the application can act on**. Speaking over the read-back withdraws the fix.
- 03 Nine tools have to be callable **mid-sentence**, with each result feeding the next thing said.

TESTS THAT SPEAK

Real audio · the Voice Agent API · the real tools

approve	the exact phrase applies the fix and the verify loop starts	<code>propose_fix → approve_fix</code>
wrong_phrase	“ yes, do it ” changes nothing — agreement is not consent	<code>propose_fix only</code>
barge_in	speaking over the read-back withdraws the proposal	<code>cancel_proposal</code>
hinglish	“ isko fix kar do ” reaches the same tool as “fix it”	<code>propose_fix</code>
contract	a Sleep Contract needs the read-back and then the phrase	<code>grant ×2</code>
undo	an applied fix can be reversed by phrase	<code>undo_fix</code>

make voice-test · each scenario speaks, then asserts on what was heard, which tools ran and how the incident ended · the right column is what must happen

MARKET

TAM

\$4.8 B

incident management software, 2026 → \$12.99 B by 2035

SAM

\$1.2 B

alert response and auto-remediation — PagerDuty alone
books \$493 M

SOM

\$125 M

small AWS teams with a one-person rotation, at list price

120,000 teams of ten or fewer × 3 responders × \$29/month · assumptions stated, argue with them

HOW IT EARNS

\$29

per responder, per month
PagerDuty is \$21–41 a seat

\$2

per **verified** remediation —
charged only when the fix was
applied *and* CloudWatch agreed
we earn nothing from noisy alerts

\$15k

per year, self-hosted
three stacks in your own account;
no data leaves

BUILT SOLO, IN TEN DAYS

SHIPPED

121 commits

302 tests

3 CloudFormation stacks

live URL, open during judging

PROVEN, NOT CLAIMED

socket killed mid-approval → **nothing ran**

barge-in → **fix withdrawn**

Hinglish audio → **tools called**

each one is a test or a harness run

NEXT

WhatsApp as the second channel

voice regression tests

contract policies for teams

a second cloud