



Beacon Night Shift

The on-call agent that fixes the 3 AM page with your voice —
and does not wake you the second time.

WHAT ON-CALL ACTUALLY COSTS

53_{min}

median time to recovery, industry-wide

12 % better in five years — against 3× the monitoring spend

74%

of DevOps engineers report burnout

on-call load is the leading indicator; seniors leave first

\$5.6k

per minute of downtime, Gartner baseline

and on a team of four there is nobody to escalate to

WHO THIS IS FOR

**A backend engineer on a team of one to five,
running production on AWS for users in another
time zone.**

No follow-the-sun rotation.
There is no second shift to
escalate to.

A small, repeating failure
surface: a rule dropped by a
deploy, a service that needs a
redploy.

No budget for an SRE
platform, and nobody to run
one.

LIVE, NOT A MOCK-UP

Speak over the read-back, and **the fix is withdrawn.**



The screenshot shows a dark green rounded rectangle representing a Beacon's interface. On the left is a circular icon with a grey outer ring, a light green inner ring, and a black center. To the right of the icon, the text reads: "Beacon is speaking — talk to interrupt", "heard: 'No, wait, stop.'", a small white square followed by "noisy room", and "Interrupt". Below this, on a white background, is a log entry. It starts with "BEACON" in bold, followed by a timestamp "00:15:24". The main text of the log entry is "A security group rule that the golden snapshot says should exist is missing, so the service can no longer —". Below the timestamp, there is a red dot followed by the word "INTERRUPTED".

Beacon is speaking — talk to interrupt
heard: "No, wait, stop."
■ noisy room
Interrupt


BEACON 00:15:24 A security group rule that the golden snapshot says should exist is missing, so the service can no longer —
● INTERRUPTED

A real run on the deployed account: the read-back cut mid-sentence, **INTERRUPTED**, `propose_fix · dry run PASSED`, and the console note — "Fix 1 withdrawn — you spoke over the read-back, so nothing was applied." The approval phrase stops working until it is proposed again.

Nine tools on one socket. Consent is decided in code.

Beacon Night Shift — the voice path

Nine tools on one socket. The model decides whether to call one; code decides whether the call is allowed.



The Voice Agent API is the mouth and the ears; every **tool.call** returns to the browser, which runs it on our Lambda with the transcript the API produced.

THE PHONE CALL

No laptop. It rings, and you answer it.

Channel 1 — you

8 kHz · G.711 μ-law



Channel 2 — Beacon

AssemblyAI · nine tools



you said “Is it fixed?”

approve_fix — the phrase is on **your** channel. Beacon read it back twice; a read-back is not consent.

CHECKED

77.6 s · get_evidence → propose_fix → approve_fix → check_recovery → check_recovery · the recording is two channels, so who approved what is a fact about the file

WHY THIS NEEDED THIS GENERATION OF MODELS

- 01 Consent has to rest on words an engineer actually said — half asleep, in two languages. A keypad cannot say which fix on which resource.
- 02 An interruption has to be an event the application can act on. Speaking over the read-back withdraws the fix.
- 03 Nine tools have to be callable mid-sentence, with each result feeding the next thing said.

TESTS THAT SPEAK

Real audio · the Voice Agent API · the real tools

approve	the exact phrase applies the fix and the verify loop starts	propose_fix → approve_fix
wrong_phrase	“yes, do it” changes nothing — agreement is not consent	propose_fix only
barge_in	speaking over the read-back withdraws the proposal	cancel_proposal
hinglish	“isko fix kar do” reaches the same tool as “fix it”	propose_fix
contract	a Sleep Contract needs the read-back and then the phrase	grant ×2
undo	an applied fix can be reversed by phrase	undo_fix
keypad	on a phone, a keypress may acknowledge — it may never approve	no consent tool

make voice-test · make phone-test · each scenario speaks, then asserts on what was heard, which tools ran and how the incident ended · the right column is what must happen

MARKET

TAM

\$4.8 B

incident management software, 2026 → \$12.99 B by 2035

SAM

\$1.2 B

alert response and auto-remediation — PagerDuty alone
books \$493 M

SOM

\$125 M

small AWS teams with a one-person rotation, at list price

120,000 teams of ten or fewer × 3 responders × \$29/month · assumptions stated, argue with them

HOW IT EARNS

\$29

per responder, per month
PagerDuty is \$21–41 a seat

\$2

per **verified** remediation —
charged only when the fix was
applied *and* CloudWatch agreed
we earn nothing from noisy alerts

\$15k

per year, self-hosted
three stacks in your own account;
no data leaves

BUILT SOLO, IN TEN DAYS

SHIPPED

141 commits

350 tests

3 channels, one socket

live URL, open during judging

PROVEN, NOT CLAIMED

socket killed mid-approval → **nothing ran**

barge-in → **fix withdrawn**

a keypress → **never consent**

each one is a test or a spoken run

NEXT

WhatsApp as the fourth channel

a carrier number, always on

contract policies for teams

a second cloud